

Deep Reinforcement Learning

Deep Q Learning

Prof. Dr. Sebastian Peitz

Chair of Safe Autonomous Systems, TU Dortmund

Summer term 2026

Content

Content

a

Where are we?

Chapter	Topic	Content
	Basics & tabular methods	
1	Multi-armed bandits	Exploration-exploitation dilemma
2	Markov decision processes	Dynamics, rewards, policies
3	Dynamic programming	Optimal decision making with <i>full knowledge</i>
4	Monte Carlo methods	<i>Data-driven learning</i> from entire episodes
5	Temporal difference learning & Q-learning	<i>Data-driven learning</i> from individual transitions
	Deep-learning-based methods	
6	Brief introduction to deep learning	The basics for what comes next
7	Value function approximation	
8	Deep Q-learning	
9	Policy gradients	
10	Actor-critic algorithms	
11	Advanced algorithms	
	Model-Based Control	

Chapter	Topic	Content
Advanced Topics		
<i>Lecture contents</i>		

Function Approximation

Motivation

- previously we were assuming that we can model everything by table loop-ups

Problems:

- real-world states can be continuous
- often we cannot see all possible states

Solution: Approximate value function, e.g. by

- linear combination of features
- decision trees
- neural networks -> DeepRL !

Value Function Approximation

- Represent state/state-action value function with a parametrized function

Linear Value Function Approximation

Weighted linear combination of features:

$$\hat{V}^{\pi}(s; \boldsymbol{\theta}) = \boldsymbol{\phi}(s)^{\top} \boldsymbol{\theta}$$

Optimization of objective (MSE):

$$J(\boldsymbol{\theta}) = \mathbb{E} \left[\left(V^{\pi}(s) - \hat{V}^{\pi}(s; \boldsymbol{\theta}) \right)^2 \right]$$

Gradient descent:

$$\Delta(\boldsymbol{\theta}) = -\frac{1}{2} \alpha \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta})$$

Update rule:

$$\Delta \boldsymbol{\theta} = -\alpha \left(V^\pi(s) - \boldsymbol{\phi}(s)^\top \boldsymbol{\theta} \right) \boldsymbol{\phi}(s)$$

Monte Carlo Value Function Approximation

Monte Carlo VFA

- G_t is a noisy but unbiased estimate of the true expected return $V^\pi(s_t)$

Update rule:

$$\Delta \theta = -\alpha (G_t - \phi(s)^\top \theta) \phi(s)$$

Monte Carlo VFA: Convergence

[Based on Marius Lindauer's lecture]

Based on Tsitsiklis and Van Roy. 1997.

Define the mean squared error of a linear value function approximation for a particular policy π relative to the true value as

$$\text{MSVE}(\boldsymbol{\theta}) = \sum_{s \in S} d(s) (V^\pi(s) - \hat{V}^\pi(s; \boldsymbol{\theta}))^2$$

where

- $d(s)$: stationary distribution of π in the true decision process
- $\hat{V}^\pi(s; \boldsymbol{\theta}) = \boldsymbol{\phi}(s)^T \boldsymbol{\theta}$, a linear value function approximation

Monte Carlo policy evaluation with VFA converges to the weights $\boldsymbol{\theta}_{MC}$ which has the minimum mean squared error possible:

$$\text{MSVE}(\boldsymbol{\theta}_{MC}) = \min_{\boldsymbol{\theta}} \sum_{s \in S} d(s) (V^{\pi}(s) - \hat{V}^{\pi}(s; \boldsymbol{\theta}))^2$$

Temporal Difference Learning with Value Function Approximation

TD-Learning with VFA

Compute target using bootstrapping:

$$r_t + \gamma \hat{V}^\pi(s_{t+1}; \boldsymbol{\theta})$$

Since target is not updated, $\hat{V}^\pi(s_{t+1}; \boldsymbol{\theta})$ is treated as a constant in the derivative.

Update rule:

$$\Delta \boldsymbol{\theta} = -\alpha \left(r_t + \gamma \hat{V}^\pi(s_{t+1}; \boldsymbol{\theta}) - \phi(s)^\top \boldsymbol{\theta} \right) \phi(s)$$

TD-Learning VFA: Convergence

[Based on Marius Lindauer's lecture]

TD(0) policy evaluation with VFA converges to weights $\boldsymbol{\theta}_{TD}$ which is a constant factor of the minimum mean squared error possible:

$$\text{MSVE}(\boldsymbol{\theta}_{TD}) \leq \frac{1}{1-\gamma} \min_{\boldsymbol{\theta}} \sum_{s \in S} d(s) (V^{\pi}(s) - \hat{V}(s; \boldsymbol{\theta}))^2$$

SARSA and Q-Learning with VFA

[Based on Marius Lindauer's lecture]

Similar to $V(s)$, we can approximate $Q(s,a)$:

$$\hat{Q}(s, a; \boldsymbol{\theta}) = \boldsymbol{\phi}(s, a)^\top \boldsymbol{\theta}$$

Monte Carlo Update:

$$\Delta \boldsymbol{\theta} = -\alpha \left(G_t - \hat{Q}(s, a; \boldsymbol{\theta}) \right) \nabla_{\boldsymbol{\theta}} \hat{Q}(s, a; \boldsymbol{\theta})$$

SARSA with TD target:


$$\Delta \boldsymbol{\theta} = -\alpha \left(r_t + \gamma \hat{Q}(s_{t+1}, a_{t+1}; \boldsymbol{\theta}) - \hat{Q}(s, a; \boldsymbol{\theta}) \right) \nabla_{\boldsymbol{\theta}} \hat{Q}(s, a; \boldsymbol{\theta})$$

Q-Learning with TD target:

$$\Delta \boldsymbol{\theta} = -\alpha \left(r_t + \gamma \max_a \hat{Q}(s_{t+1}, a; \boldsymbol{\theta}) - \hat{Q}(s, a; \boldsymbol{\theta}) \right) \nabla_{\boldsymbol{\theta}} \hat{Q}(s, a; \boldsymbol{\theta})$$

Deep Q-Learning

Towards more complex learning tasks



CartPole:

Atari:



Credit: Mnih et al. 2013

-> Deep Neural Networks (DNNs) as value function approximators

Recall: Online Q-Learning with VFA

1. Take action a_t and observe (s_t, a_t, r_t, s_{t_1}) -> correlated! breaks i.i.d assumption of NNs
2. Update Q-Network:

$$\Delta\theta = -\alpha \left(r_t + \gamma \max_a \hat{Q}(s_{t+1}, a; \theta) - \underbrace{\hat{Q}(s, a; \theta)}_{\text{non-stationary!}} \right) \nabla_{\theta} \hat{Q}(s, a; \theta)$$

Deep Q-Network (DQN)

1. Correlation -> **Replay Buffer** from which we sample batches i.i.d.
2. **Target network**: Copy of previous network weights used as target and updated delayed: θ^-

Maximization Bias

- maximum in both value estimation and action selection
- leads to positive bias in Q-values

Double DQN

- keep two q networks and toss coin which one to update

Prioritized Experience Replay

Based on Marius Lindauer's Lecture

- Let i be the index of the i -th tuple of experience (s_i, a_i, r_i, s_{i+1})
- Sample tuples for the update using priority function
- Priority of a tuple i is proportional to DQN error

$$p_i = |r_i + \gamma \max_{a' \in \mathcal{A}} Q(s_{i+1}, a'; \theta^-) - Q(s_i, a_i; \theta)|$$

- Update p_i every update. p_i for new tuples is set to the maximum value

- One method: proportional (stochastic prioritization)

$$P(i) = \frac{p_i^\beta}{\sum_k p_k^\beta}$$

- $\beta = 0$ yields random selections

<https://arxiv.org/pdf/1511.05952>

References