

# Deep Reinforcement Learning

## *Advanced Algorithms (Part I)*

Prof. Dr. Sebastian Peitz

Chair of Safe Autonomous Systems, TU Dortmund

Summer term 2026

# Content

# Content

- The evolution of modern RL algorithms
- Part (I): Improving policy gradient methods
  - Natural policy gradient
  - Trust-region policy optimization (TRPO)
  - Proximal policy optimization (PPO)

# Where are we?

| Chapter | Topic                                                                  | Content                                       |
|---------|------------------------------------------------------------------------|-----------------------------------------------|
|         | <b>Basics &amp; tabular methods</b>                                    |                                               |
| 1-5     | Bandits, MDPs, Dynamic Programming, Monte Carlo, TD Learning           | RL basics in finite dimensions                |
|         | <b>Deep-learning-based methods</b>                                     |                                               |
| 6       | Brief introduction to deep learning                                    | The basics for what comes next                |
| 7       | Value function approximation                                           | Value estimation with function approximation  |
| 8       | Deep Q-learning                                                        | Q-learning with neural networks               |
| 9       | Policy gradients                                                       | Direct optimization of the policy             |
| 10      | Actor-critic algorithms                                                | Improved policy gradients via value functions |
| 11      | Advanced algorithms (Part I): From policy gradient to PPO              | The evolution of modern RL algorithms         |
| 12      | Advanced algorithms (Part II): From $Q$ -learning to Soft Actor-Critic |                                               |
| 13      | Exploration                                                            |                                               |
|         | <b>Model-Based Control</b>                                             |                                               |
|         | <b>Advanced Topics</b>                                                 |                                               |

| Chapter | Topic | Content |
|---------|-------|---------|
|---------|-------|---------|

*Lecture contents*

# The evolution of modern RL algorithms

# The evolution of modern RL algorithms

## (I) Policy gradient methods

- Policy is explicitly parameterized and optimized directly:

$$L_{\pi}(\phi) = \mathbb{E}_{\tau \sim p_{\phi}(\tau)} [r(\tau)].$$

- Gradient ascent:  $\phi \leftarrow \phi + \alpha \nabla_{\phi} L_{\pi}(\phi)$ .
- Methods are typically **on-policy**.
- **Algorithms:** REINFORCE  $\rightarrow$  Actor-Critic  $\rightarrow$  Natural Policy Gradient  $\rightarrow$  Trust-region policy optimization (TRPO)  $\rightarrow$  Proximal policy optimization (PPO).

### Shortcomings

- High variance gradients.
- Unstable policy updates.
- Catastrophic performance collapse.

### Improvement strategy

- How to safely update policies.

## (II) Value-based methods

- Approximation of the  $Q$ -function:

$$Q^*(s, a) = r + \max_{a' \in \mathcal{A}} Q^*(s', a').$$

- Extract implicit policy:  $\pi(s) = \arg \max_{a \in \mathcal{A}} Q^*(s, a)$ .
- Typically **off-policy**.
- **Algorithms:** Q-learning  $\rightarrow$  DQN  $\rightarrow$  Deep deterministic policy gradient (DDPG)  $\rightarrow$  TD3  $\rightarrow$  Soft actor-critic (SAC).

### Shortcomings

- Instability from bootstrapping.
- Overestimation bias.
- Maximization over actions / continuous action spaces.


### Improvement strategy

- Stabilizing  $Q$ -learning with function approximation.

- Solving the continuous argmax problem.

# Part (I): Improving policy gradient methods

# Ill-conditioned gradients




A simple exemplary MDP (Peters and Schaal 2008). Inspired by Sergey Levine's CS285 lecture.

$$r_t = -s_t^2 - a_t^2$$

$$\log \pi_\phi(a_t|s_t) = -\frac{1}{2\sigma^2}(ks_t - a_t)^2 + C, \quad \phi = (k, \sigma).$$


$$\nabla_\phi \log \pi_\phi(a_t|s_t) = \begin{pmatrix} -\frac{(ks_t - a_t)s_t}{\sigma^2} \\ \frac{(ks_t - a_t)^2}{\sigma^3} \end{pmatrix}$$

- **Optimum:**  $\phi^* = (-1, 0)$ .
- **Issue:** The gradients do not point towards the optimum for smaller  $\sigma$ !
  - The  $\sigma^{-3}$  term blows up.



Source: (Peters and Schaal 2008)

Closely related to ill-conditioned optimization problems:





# Constraining the ascent direction

- Recall the policy gradient update:  $\phi \leftarrow \phi + \alpha \nabla_{\phi} L_{\pi}(\phi)$ .
- A natural idea: **Constrain the length of our update step!**
  - Linearization of our optimization problem around  $\phi$  using **Taylor series expansion**:


$$\begin{aligned} L(\phi') &= L_{\pi}(\phi) + \sum_{i=1}^n \frac{\partial L}{\partial \phi_i} (\phi'_i - \phi_i) + \mathcal{O}((\phi' - \phi)^2) \\ &= L_{\pi}(\phi) + (\phi' - \phi)^{\top} \nabla_{\phi} L_{\pi}(\phi). \end{aligned}$$

- Constrained optimization of  $\phi'$  along the steepest ascent direction:

$$\phi' \leftarrow \arg \max_{\phi'} (\phi' - \phi)^{\top} \nabla_{\phi} L_{\pi}(\phi) \quad \text{subject to} \quad \|\phi' - \phi\|^2 \leq \epsilon.$$

- But:** We just saw that some parameters change probabilities *a lot more than others!*
- Alternative:** Rescale the gradient so that this does not happen!
- A better way to constrain the update distance: The **change in distribution** via the **Kullback-Leibler (KL) divergence!**

$$D_{\text{KL}}(\pi_{\phi} \parallel \pi'_{\phi}).$$



Source: (Peters and Schaal 2008)


## Aside: KL-divergence and Fisher information matrix

# Kullback-Leibler divergence

- The **Kullback-Leibler divergence** (or **KL divergence**) measures the “distance” between two distributions (e.g.,  $p(x)$  and  $q(x)$ )

$$D_{\text{KL}}(p\|q) = \begin{cases} \sum_{x \in \mathcal{X}} p(x) \log\left(\frac{p(x)}{q(x)}\right) & \text{if } \mathcal{X} \text{ is finite} \\ \int_{\mathcal{X}} p(x) \log\left(\frac{p(x)}{q(x)}\right) dx & \text{if } \mathcal{X} \text{ is continuous} \end{cases}$$

- It is zero if and only if  $p(x) = q(x)$ .
- **Intuition:** Imagine you have two different maps of the exact same city.
  - Map A is completely accurate. Every street, etc., is exactly where it should be.
  - Map B was drawn from memory by a friend. It's mostly right with some minor flaws.
  - If you use Map B to navigate the city, you are going to make some mistakes.
  - The KL divergence answers the question: “How much extra ‘gas’ (or information) will I waste if I use Map B instead of Map A?”
- KL Divergence measures the *surprise* or *unfair penalty* you get when you use  $q$  to predict something that actually follows  $p$ . If  $q$  is a perfect match for  $p$  (i.e.,  $q = p$ ), your surprise is zero. The worse your guess ( $q$ ) is, the higher the KL divergence.
- **Question:** What happens if our “model”  $q$  is absolutely certain that something is **not** going to happen ( $q(x) = 0$  for some  $x$ )?
  - 💡 The KL divergence is not a real distance. For instance, it is not symmetric (i.e.,  $D_{\text{KL}}(p\|q) \neq D_{\text{KL}}(q\|p)$  in general), which means that it also does not satisfy the triangle inequality. Still, it gives us a measure of how far two distributions are apart.



[Source].

# Fisher information matrix

- The **Fisher information matrix (FIM)** measures how much information an observable random variable  $x$  carries about an unknown parameter  $\phi$  (e.g., the weights of a neural network).
- Mathematically, if we have a log-likelihood function  $\log p(x|\phi)$ , the FIM (denoted as  $F$ ) is the variance of its gradient (the “score”):

$$F = \mathbb{E} \left[ (\nabla_{\phi} \log p(x|\phi)) (\nabla_{\phi} \log p(x|\phi))^{\top} \right]. \quad (1)$$

**In simpler terms:** How do variations of  $\phi$  influence the resulting probability distribution?

- High Fisher Information: Small change in  $\phi \Rightarrow$  large shift in the distribution.
- Low Fisher Information: Change in  $\phi \Rightarrow$  only small shift in the distribution.

## Relation to the KL divergence

- Consider a very small change in  $\phi$  by  $\Delta\phi$ , and study the KL divergence  $D_{\text{KL}}(p(x|\phi) \| p(x|\phi + \Delta\phi))$ .
- If we perform a **Taylor series expansion** of  $D_{\text{KL}}$  around  $\Delta\phi = 0$ , then the first two terms are zero ( $D_{\text{KL}}(p(x|\phi) \| p(x|\phi)) = 0$  and since  $\Delta\phi = 0$  is the minimum, the first derivative is zero as well).
- The second-order term is thus the leading term!  $\Rightarrow$  for small  $\Delta\phi$ , we have

$$D_{\text{KL}}(p(x|\phi) \| p(x|\phi + \Delta\phi)) \approx \frac{1}{2} \Delta\phi^{\top} F \Delta\phi. \quad (2)$$

# KL-divergence and Fisher information in RL

- When using the KL divergence or the Fisher information matrix in RL, we often do so in terms of different policies.
- This means that we consider the difference in the **action distribution**  $D_{\text{KL}}(\pi_\phi \parallel \pi_{\phi'})$ .
- However, this expression is incomplete or even misleading.
- Instead, we consider
  - for a fixed  $s$ ,

$$D_{\text{KL}}(\pi_\phi(\cdot|s) \parallel \pi_{\phi'}(\cdot|s)) = \int_{\mathcal{A}} \pi_\phi(a|s) \log \left( \frac{\pi_\phi(a|s)}{\pi_{\phi'}(a|s)} \right) da,$$

- the expectation over the state space, i.e.,

$$\bar{D}_{\text{KL}}(\pi_\phi \parallel \pi_{\phi'}) = \mathbb{E}_{s \sim p_\phi} [D_{\text{KL}}(\pi_\phi(\cdot|s) \parallel \pi_{\phi'}(\cdot|s))],$$

- or the maximum entry, i.e.,

$$D_{\text{KL}}^{\max}(\pi_\phi \parallel \pi_{\phi'}) = \max_{s \in \mathcal{S}} D_{\text{KL}}(\pi_\phi(\cdot|s) \parallel \pi_{\phi'}(\cdot|s)).$$

- The relation between KL divergence and FIM for small changes can be derived analogously to the previous slide and reads

$$\mathbb{E}_{s \sim p_\phi} [D_{\text{KL}} (\pi_\phi \| \pi_{\phi'})] \approx \frac{1}{2} \Delta \phi^\top \left( \underbrace{\mathbb{E}_{s \sim p_\phi} [F(s)]}_{=F} \right) \Delta \phi = \frac{1}{2} \Delta \phi^\top F \Delta \phi.$$

# Natural policy gradient

# Back to constrained ascent directions

- Recall: constraining w.r.t. the parameter  $\phi$  does not solve the issue of strongly different impacts of individual parameters:

$$\phi' \leftarrow \arg \max_{\phi'} (\phi' - \phi)^\top \nabla_\phi L_\pi(\phi) \quad \text{subject to} \quad \|\phi' - \phi\|^2 \leq \epsilon.$$

- Question:** What is a good alternative to avoid too large (and thus, harmful) steps?  
 $\Rightarrow$  the KL divergence!

$$\phi' \leftarrow \arg \max_{\phi'} (\phi' - \phi)^\top \nabla_\phi L_\pi(\phi) \quad \text{subject to} \quad \mathbb{E}_{s \sim p_\phi} [D_{\text{KL}}(\pi_{\phi'}(\cdot|s) \parallel \pi_\phi(\cdot|s))] \leq \epsilon.$$

- Constraining the KL divergence means that we automatically treat different parameters differently!
- If a change in one of the entries of  $\phi$  results in large policy changes, then the constraint does not allow large changes in that entry.
- Since we want to have small updates (constrained by  $\epsilon$ ): approximation via the Fisher information matrix,

$$\phi' \leftarrow \arg \max_{\phi'} (\phi' - \phi)^\top \nabla_\phi L_\pi(\phi) \quad \text{subject to} \quad (\phi' - \phi)^\top F (\phi' - \phi) \leq \epsilon,$$

with  $F = \mathbb{E}_{s \sim p_\phi, a \sim \pi_\phi(\cdot|s)} [(\nabla_\phi \log \pi_\phi(a|s))(\nabla_\phi \log \pi_\phi(a|s))^\top].$

# Natural policy gradient



- We have now defined a policy gradient whose update length is constrained in terms of the KL divergence:

$$\phi' \leftarrow \arg \max_{\phi'} (\phi' - \phi)^\top \nabla_{\phi} L_{\pi}(\phi) \quad \text{subject to} \quad (\phi' - \phi)^\top F(\phi' - \phi) \leq \epsilon. \quad (3)$$

- What remains is the question of solving (3).
- Without going into further details (e.g., (S. M. Kakade 2001; Peters and Schaal 2008))  
 $\Rightarrow$  Solution: **scale** the policy gradient **by the inverse Fisher information matrix**,

$$\phi \leftarrow \phi + \alpha F^{-1} \nabla_{\phi} L_{\pi}(\phi). \quad (4)$$

- **Intuition** behind the inverse FIM  $F^{-1}$ :
  - Parameters with a high impact have high Fisher information.
  - Scaling by the inverse “normalizes” the individual impacts.
- This technique is known as the **natural policy gradient** (also **covariant policy gradient**).
  - Preconditioning steps of this type are very common in optimization in general.
  - Drawback:  $F \in \mathbb{R}^{d \times d}$  can be very expensive to calculate (and invert) for very large parameter vectors.



Source: (Peters and Schaal 2008)

# Limitations of the natural policy gradient

$$\phi \leftarrow \phi + \alpha F^{-1} \nabla_{\phi} L_{\pi}(\phi) \quad (4)$$

## 1. The step size $\alpha$ is hard to choose!

- Scaling doesn't prevent the step from changing the policy too much.
- Large steps can collapse performance.

## 2. No explicit control of policy change.

- NPG does not explicitly constrain the distance between  $\pi_{\text{old}}$  and  $\pi_{\text{new}}$ .
- Data for gradient estimation is sampled from  $\pi_{\text{old}}$ .
- If new policy differs too much, gradient estimate becomes unreliable.

## 3. Inversion of the Fisher matrix.

- NPG requires inversion of the Fisher Information Matrix.
- If our network has  $d$  parameters, the FIM is a  $d \times d$  matrix.
- For a modest deep learning network with 100,000 parameters,  $F$  contains 10 billion elements.

## 4. Performance guarantees require small policy changes.

- NPG does not provide a mathematical guarantee that  $\pi_{\text{new}}$  will actually be better than (or at least as good as)  $\pi_{\text{old}}$ .
- Theory shows that small policy updates should improve performance (S. Kakade and Langford 2002) ...
- ... but the basic NPG update does not enforce the required small change condition.

- Storing this in memory and computing its exact inverse ( $\mathcal{O}(d^3)$  complexity)  $\Rightarrow$  computationally impossible.

# Trust-region policy optimization

# Performance measures of policies

- Let's introduce a new (and yet *already known* quantity): the **performance measure** of a policy,

$$\eta(\pi) = \mathbb{E}_{\tau \sim \pi} \left[ \sum_{t=0}^{T-1} \gamma^t r_t \right] = \mathbb{E}_{s_0 \sim p_0} [V^\pi(s_0)].$$

- Remember our reinforcement learning objective

$$L_\pi(\phi) = \mathbb{E}_{\tau \sim p_\phi(\tau)} \left[ \sum_{t=0}^{T-1} r_t \right] = \mathbb{E}_{s_0 \sim p_0} [V^{\pi_\phi}(s_0)]?$$

- If we consider a policy parameterized by  $\phi$ , i.e.,  $\pi_\phi$ , then the two equations are the same,

$$\eta(\pi_\phi) = L_\pi(\phi).$$

- The difference is more in terms of the *motivation*, not the definition:
  - We use  $L_\pi(\phi)$  in policy gradients, since we want to find ascent directions for the network parameters.
  - We use  $\eta(\pi)$  as a global quantity to establish mathematical lower-bounds so that an algorithm can compute a safe step size (the trust region).

# Comparing the performance measure of two policies (1)

- Can we relate two policies  $\pi_{\text{old}}$  and  $\pi_{\text{new}}$  using advantage functions?
- Let's use  $\gamma = 1$  for simplicity (but one can work this out with  $\gamma$  as well).
- Consider the advantage function under the old policy  $\pi_{\text{old}}$ ,

$$A_{\pi_{\text{old}}}(s_t, a_t) = r_t + V^{\pi_{\text{old}}}(s_{t+1}) - V^{\pi_{\text{old}}}(s_t).$$

- Now sample a trajectory under the new policy  $\pi_{\text{new}}$  (i.e.,  $(s_0, a_0) \rightarrow \dots \rightarrow (s_{T-1}, a_{T-1})$ ) and calculate the advantages:

$$\begin{aligned} A_{\pi_{\text{old}}}(s_0, a_0) &= r_0 + V^{\pi_{\text{old}}}(s_1) - V^{\pi_{\text{old}}}(s_0) \\ A_{\pi_{\text{old}}}(s_1, a_1) &= r_1 + V^{\pi_{\text{old}}}(s_2) - V^{\pi_{\text{old}}}(s_1) \\ &\vdots \\ A_{\pi_{\text{old}}}(s_{T-2}, a_{T-2}) &= r_{T-2} + V^{\pi_{\text{old}}}(s_{T-1}) - V^{\pi_{\text{old}}}(s_{T-2}) \\ A_{\pi_{\text{old}}}(s_{T-1}, a_{T-1}) &= r_{T-1} + \underbrace{V^{\pi_{\text{old}}}(s_T) - V^{\pi_{\text{old}}}(s_{T-1})}_{=0 \text{ (terminal)}} \end{aligned}$$

# Comparing the performance measure of two policies (2)

- Take the sum:

$$\begin{aligned} \sum_{t=0}^{T-1} A_{\pi_{\text{old}}}(s_t, a_t) &= \underbrace{r_0 + V^{\pi_{\text{old}}}(s_1) - V^{\pi_{\text{old}}}(s_0)}_{=A_{\pi_{\text{old}}}(s_0, a_0)} + \underbrace{r_1 + V^{\pi_{\text{old}}}(s_2) - V^{\pi_{\text{old}}}(s_1)}_{=A_{\pi_{\text{old}}}(s_1, a_1)} + \dots + \underbrace{r_{T-2} + V^{\pi_{\text{old}}}(s_{T-1}) - V^{\pi_{\text{old}}}(s_{T-2})}_{=A_{\pi_{\text{old}}}(s_{T-2}, a_{T-2})} + \underbrace{r_{T-1} - V^{\pi_{\text{old}}}(s_{T-1})}_{=A_{\pi_{\text{old}}}(s_{T-1}, a_{T-1})} \\ &= r_0 + r_1 + \dots + r_{T-2} + r_{T-1} - V^{\pi_{\text{old}}}(s_0) = \left( \sum_{t=0}^{T-1} r_t \right) - V^{\pi_{\text{old}}}(s_0). \end{aligned}$$

- Shift around and take the expectation with respect to the new policy (in the case of  $V^{\pi_{\text{old}}}$ , this reduces to  $s_0 \sim p_0$ ):

$$\begin{aligned} \mathbb{E}_{s \sim p_{\pi_{\text{new}}}, a \sim \pi_{\text{new}}} \left[ \left( \sum_{t=0}^{T-1} r_t \right) \right] &= \mathbb{E}_{s_0 \sim p_0} [V^{\pi_{\text{old}}}(s_0)] + \mathbb{E}_{s \sim p_{\pi_{\text{new}}}, a \sim \pi_{\text{new}}} \left[ \sum_{t=0}^{T-1} A_{\pi_{\text{old}}}(s_t, a_t) \right] \\ \Leftrightarrow \quad \eta(\pi_{\text{new}}) &= \eta(\pi_{\text{old}}) + \mathbb{E}_{s \sim p_{\pi_{\text{new}}}, a \sim \pi_{\text{new}}} \left[ \sum_{t=0}^{T-1} \gamma^t A_{\pi_{\text{old}}}(s_t, a_t) \right] \end{aligned}$$

- Right now, the expectation is written over an entire lifetime trajectory ( $\tau \sim \pi_{\text{new}}$ ).
- To be useful for an optimization algorithm, we need to break it down step-by-step and look at individual states and actions.



# Comparing the performance measure of two policies (3)

- Let's consider the continual case ( $T = \infty$ ) such that we can dissect the expectation into individual steps:

$$\begin{aligned}\mathbb{E}_{s \sim p_{\pi_{\text{new}}}, a \sim \pi_{\text{new}}} \left[ \sum_{t=0}^{\infty} \gamma^t A_{\pi_{\text{old}}}(s_t, a_t) \right] &= \sum_{t=0}^{\infty} \gamma^t \mathbb{E}_{s \sim p_{\pi_{\text{new}}}, a \sim \pi_{\text{new}}} [A_{\pi_{\text{old}}}(s_t, a_t)] \\ &= \sum_{t=0}^{\infty} \int_{\mathcal{S}} p_{\pi_{\text{new}}}(s) \int_{\mathcal{A}} \pi_{\text{new}}(a|s) A_{\pi_{\text{old}}}(s, a) da ds \\ &= \int_{\mathcal{S}} \underbrace{\left( \sum_{t=0}^{\infty} \gamma^t p_{\pi_{\text{new}}}(s) \right)}_{=\rho_{\pi_{\text{new}}}(s)} \left( \int_{\mathcal{A}} \pi_{\text{new}}(a|s) A_{\pi_{\text{old}}}(s, a) da \right) ds \\ &= \mathbb{E}_{s \sim \rho_{\pi_{\text{new}}}, a \sim \pi_{\text{new}}} [A_{\pi_{\text{old}}}(s, a)].\end{aligned}$$

- Here, we have introduced the **discounted state distribution**  $\rho_{\pi_{\text{new}}}(s)$ .

## Relation between the performance measures

$$\eta(\pi_{\text{new}}) = \eta(\pi_{\text{old}}) + \mathbb{E}_{s \sim \rho_{\pi_{\text{new}}}, a \sim \pi_{\text{new}}} [A_{\pi_{\text{old}}}(s, a)] \quad (5)$$

💡 When sampling, the data will automatically be distributed according to  $\rho_{\pi_{\text{new}}}(s)$ . See also the discussion in the actor-critic lecture.



# Comparing the performance measure of two policies (4)

Relation between the performance measures

$$\eta(\pi_{\text{new}}) = \eta(\pi_{\text{old}}) + \mathbb{E}_{s \sim \rho_{\pi_{\text{new}}}, a \sim \pi_{\text{new}}} [A_{\pi_{\text{old}}}(s, a)] \quad (5)$$

What's so surprising about this equation?

- The Advantage function is built entirely out of the old policy's values,
- yet when we weight it by the new policy's state distribution,
- it perfectly calculates the new policy's total return.

Upside and downside of this identity

- + Proof that if you can find a new policy with a positive expected advantage under the old policy, your policy will get better.
- To compute the exact value, we must sample states from the new policy ( $s \sim p_{\pi_{\text{new}}}$ ), which we do not yet know!

**Solution approach in TRPO:** “If  $\pi_{\text{new}}$  is very close to  $\pi_{\text{old}}$  (enforced by a KL-constraint), we can swap the state distribution  $p_{\pi_{\text{new}}}$  for  $p_{\pi_{\text{old}}}$  and safely optimize an approximation (i.e., a *surrogate objective*).”

# How to turn this into an algorithm

## Relation between the performance measures

$$\eta(\pi_{\text{new}}) = \eta(\pi_{\text{old}}) + \mathbb{E}_{s \sim \rho_{\pi_{\text{new}}}, a \sim \pi_{\text{new}}} [A_{\pi_{\text{old}}}(s, a)] \quad (5)$$

Starting with (5), we need to take the several steps to arrive at an algorithm that we can actually implement.

1. Construct a local approximation of (5) (the *surrogate loss*) which we can estimate using samples.
2. Ensure that this approximation is sufficiently accurate.
3. Turn the surrogate loss function into a practical optimization problem.
4. Make sure that it scales to large problems.

## Trust-region policy optimization (TRPO)

The first realization of these steps resulted in the TRPO algorithm, which was derived in (Schulman et al. 2015).

# TRPO: The surrogate loss function

## Relation between the performance measures

$$\eta(\pi_{\text{new}}) = \eta(\pi_{\text{old}}) + \mathbb{E}_{s \sim \rho_{\pi_{\text{new}}}, a \sim \pi_{\text{new}}} [A_{\pi_{\text{old}}}(s, a)] \quad (5)$$

- In principle, this equation shows us a way to guarantee policy improvement.
- If we can ensure positive expected advantages, then we improve!
- **But:** Sampling under the new policy ( $s \sim \rho_{\pi_{\text{new}}}, a \sim \pi_{\text{new}}$ ) is infeasible, as we do not have it yet!

## A surrogate loss function

- Simple trick: just sample the state  $s$  under the *old* policy!

$$L_{\pi_{\text{old}}}(\pi) = \eta(\pi_{\text{old}}) + \mathbb{E}_{s \sim \rho_{\pi_{\text{old}}}, a \sim \pi} [A_{\pi_{\text{old}}}(s, a)]. \quad (6)$$

- Under the assumption that  $\pi_{\text{old}}$  and  $\pi_{\text{new}}$  do not differ too much, this is a reasonable assumption.
- In the limit case  $\pi_{\text{old}} = \pi_{\text{new}}$ , we have equality of (5) and (6), as well as their gradients (Schulman et al. 2015, Eq. (4)).

# TRPO: Approximation error

In (Schulman et al. 2015), it was shown that error between the two expressions (5) and (6) can be bounded

$$\eta(\pi) \geq L_{\pi_{\text{old}}}(\pi) - C \cdot D_{\text{KL}}^{\max}(\pi_{\text{old}} \parallel \pi), \quad \text{with} \quad D_{\text{KL}}^{\max}(\pi_{\text{old}} \parallel \pi) = \max_{s \in \mathcal{S}} D_{\text{KL}}(\pi_{\phi}(\cdot | s) \parallel \pi_{\phi'}(\cdot | s)),$$

and  $C$  is a constant derived from the environment's transition dynamics and the maximum bounds of the advantage function.

**Central message:** If we maximize the right-hand side of this inequality, we are guaranteed to improve (or at least maintain) our true performance  $\eta(\pi)$ !

**Maximizing the surrogate loss**  $\Rightarrow$  introducing a parametric policy  $\pi_{\phi}$ , our surrogate-based optimization problem is

$$\max_{\phi} L_{\pi_{\text{old}}}(\pi_{\phi}) - C \cdot D_{\text{KL}}^{\max}(\pi_{\text{old}} \parallel \pi_{\phi}),$$

or, if we consider two iterates  $\phi$  and  $\phi'$

$$\max_{\phi'} L_{\pi_{\phi}}(\pi_{\phi'}) - C \cdot D_{\text{KL}}^{\max}(\pi_{\phi} \parallel \pi_{\phi'}).$$

## Challenges:

1. The constant  $C$  is usually large and thus restrictive  
**Solution:** Transfer to a constraint with hyperparameter  $\delta$ .

2. Computing the maximum KL divergence over all possible states ( $D_{\text{KL}}^{\max}(\pi_{\phi_{\text{old}}} \parallel \pi_{\phi})$ ) is impossible.  
**Solution:** Replace by the average KL divergence over the states that were actually observed:

$$\bar{D}_{\text{KL}}(\pi_{\phi} \parallel \pi_{\phi_{\text{old}}}) = \mathbb{E}_{s \sim \rho_{\pi_{\phi_{\text{old}}}}} [D_{\text{KL}}(\pi_{\phi_{\text{old}}}(\cdot | s) \parallel \pi_{\phi}(\cdot | s))].$$

# TRPO: Sampling the correct expectation

Our final observation is that in our surrogate loss function  $L_{\pi_{\text{old}}}$  (Eq. (6)), the “old” performance measure  $\eta(\pi_{\text{old}})$  serves as a constant  $\Rightarrow$  we can neglect it!

- Overall, we obtain the following surrogate optimization problem:

$$\max_{\phi} \underbrace{\mathbb{E}_{s \sim \rho_{\pi_{\text{old}}}, a \sim \pi_{\phi}} [A_{\pi_{\text{old}}}(s, a)]}_{= L_{\pi_{\text{old}}}(\pi_{\phi}) - \eta(\pi_{\text{old}})} \quad \text{subject to} \quad \underbrace{\mathbb{E}_{s \sim \rho_{\pi_{\text{old}}}} [D_{\text{KL}}(\pi_{\text{old}}(\cdot|s) \| \pi_{\phi}(\cdot|s))]}_{= \bar{D}_{\text{KL}}(\pi_{\text{old}} \| \pi_{\phi})} \leq \delta.$$

- What’s wrong with this equation?

$\Rightarrow$  In the loss, we’re compute expectations over the new policy  $\pi_{\phi}$ , but the data was collected under  $\pi_{\text{old}}$ !

$$\mathbb{E}_{s \sim \rho_{\pi_{\text{old}}}, a \sim \pi_{\phi}} [A_{\pi_{\text{old}}}(s, a)] = \int_{\mathcal{S}} \rho_{\pi_{\text{old}}}(s) \int_{\mathcal{A}} \pi_{\phi}(a|s) A_{\pi_{\text{old}}}(s, a) \, da \, ds.$$

- Solution:** Importance sampling!

$$\int_{\mathcal{S}} \rho_{\pi_{\text{old}}}(s) \int_{\mathcal{A}} \pi_{\text{old}}(a|s) \frac{\pi_{\phi}(a|s)}{\pi_{\text{old}}(a|s)} A_{\pi_{\text{old}}}(s, a) \, da \, ds = \mathbb{E}_{s \sim \rho_{\pi_{\text{old}}}, a \sim \pi_{\text{old}}} \left[ \frac{\pi_{\phi}(a|s)}{\pi_{\text{old}}(a|s)} A_{\pi_{\text{old}}}(s, a) \right].$$

# TRPO: The surrogate optimization problem

Trust-region policy optimization (TRPO) optimization problem

$$\max_{\phi} \underbrace{\mathbb{E}_{s \sim \rho_{\pi_{\text{old}}}, a \sim \pi_{\text{old}}} \left[ \frac{\pi_{\phi}(a|s)}{\pi_{\text{old}}(a|s)} A_{\pi_{\text{old}}}(s, a) \right]}_{=L_{\text{TRPO}}(\phi)} \quad \text{subject to} \quad \underbrace{\mathbb{E}_{s \sim \rho_{\pi_{\text{old}}}} [D_{\text{KL}}(\pi_{\text{old}}(\cdot|s) \| \pi_{\phi}(\cdot|s))]}_{=\bar{D}_{\text{KL}}(\pi_{\text{old}} \| \pi_{\phi})} \leq \delta. \quad (7)$$

Next question: How do we solve it efficiently?  $\Rightarrow$  Let's start with the gradient!

$$\nabla_{\phi} L_{\text{TRPO}}(\phi) = \nabla_{\phi} \mathbb{E}_{s \sim \rho_{\pi_{\text{old}}}, a \sim \pi_{\text{old}}} \left[ \frac{\pi_{\phi}(a|s)}{\pi_{\text{old}}(a|s)} A_{\pi_{\text{old}}}(s, a) \right] = \mathbb{E}_{s \sim \rho_{\pi_{\text{old}}}, a \sim \pi_{\text{old}}} \left[ \frac{\nabla_{\phi} \pi_{\phi}(a|s)}{\pi_{\text{old}}(a|s)} A_{\pi_{\text{old}}}(s, a) \right]$$

The Policy gradient for the TRPO surrogate loss

$$\begin{aligned}
g &= \nabla_{\phi} L_{\text{TRPO}}(\phi) \Big|_{\phi=\phi_{\text{old}}} = \mathbb{E}_{s \sim \rho_{\pi_{\phi_{\text{old}}}}, a \sim \pi_{\phi_{\text{old}}}} \left[ \frac{\nabla_{\phi} \pi_{\phi}(a|s) \Big|_{\phi=\phi_{\text{old}}}}{\pi_{\phi_{\text{old}}}(a|s)} A_{\pi_{\phi_{\text{old}}}}(s, a) \right] \\
&= \mathbb{E}_{s \sim \rho_{\pi_{\phi_{\text{old}}}}, a \sim \pi_{\phi_{\text{old}}}} \left[ \nabla_{\phi} \log \pi_{\phi}(a|s) \Big|_{\phi=\phi_{\text{old}}} A_{\pi_{\phi_{\text{old}}}}(s, a) \right] \quad (\text{Log-identity trick})
\end{aligned} \tag{8}$$

is the standard policy gradient with value baseline!

# TRPO: Solving the optimization problem (7)

To solve Problem (7) efficiently, we need two steps:

1. Linearize the objective function: First-order **Taylor series expansion** around  $\phi_{\text{old}}$  using the gradient  $g$  (Eq. (8)).

$$L_{\text{TRPO}}(\phi) \approx L_{\text{TRPO}}(\phi_{\text{old}}) + g^\top (\phi - \phi_{\text{old}}) = g^\top (\phi - \phi_{\text{old}}).$$

(The term  $L_{\text{TRPO}}(\phi_{\text{old}})$  drops out because it is exactly zero at  $\phi_{\text{old}}$ : No advantage and sampling ratio is one.)

2. Quadratic approximation of the KL constraint for small  $\Delta\phi = (\phi - \phi_{\text{old}}) \Rightarrow$  Fisher information matrix (Eq. (2))!

## Linear-quadratic approximation of the TRPO problem (7)

$$\max_{\Delta\phi} g^\top \Delta\phi \quad \text{subject to} \quad \frac{1}{2} \Delta\phi^\top F \Delta\phi \leq \delta, \quad (9)$$

where, according to (1),

$$F(\phi) = \mathbb{E}_{s \sim \rho_{\phi_{\text{old}}}, a \sim \pi_{\phi_{\text{old}}}} \left[ \left[ (\nabla_{\phi} \log \pi_{\phi}(a|s)) (\nabla_{\phi} \log \pi_{\phi}(a|s))^\top \right]_{\phi=\phi_{\text{old}}} \right].$$

## Solution to (9)

The solution (via **Lagrange multipliers**) yields

$$\Delta\phi = \sqrt{\frac{2\delta}{g^\top F^{-1} g}} F^{-1} g. \quad (10)$$

**Note:** This is NPG (4) with automatic step size selection  $\alpha = \sqrt{\frac{2\delta}{g^\top F^{-1} g}}!$

💡 In classical optimization, the local approximation (9) is solved repeatedly. The solution can be shown to converge to the optimum of (7), which is known as **Sequential Quadratic Programming (SQP)**. In TRPO, we solve the approximation (9) just once before collecting new data.

# TRPO: Avoiding matrix inversion

$$\Delta\phi = \sqrt{\frac{2\delta}{g^\top F^{-1}g}} F^{-1}g \quad (10)$$

- Matrix inversion scales cubically. That is, for a  $d$ -dimensional parameter  $\phi \in \mathbb{R}^d$ , we have  $\mathcal{O}(d^3)$ .
- Solution: Use the **conjugate gradient (CG) method** for solving linear systems of the form  $Fx = g$ .
  - Usually, CG scales according to  $\mathcal{O}(d^2)$  per iteration.
  - However, in TRPO, we can avoid calculating  $F \in \mathbb{R}^{d \times d}$  entirely using automatic differentiation  $\Rightarrow$  scales linearly with the number of weights (i.e.,  $\mathcal{O}(d)$ ).
  - This is as fast as it gets!

## Approach

- Compute  $x = F^{-1}g$  using the CG method ( $\mathcal{O}(d)$  if we use autodiff).
- Calculate the denominator in (10):  $g^\top F^{-1}g = g^\top x$ .
- Compute the step size  $\beta = \sqrt{\frac{2\delta}{g^\top x}}$ .
- Compute the update  $\Delta\phi = \beta x$ .

# TRPO: Line search

## TRPO surrogate problem

$$\max_{\phi} L_{\text{TRPO}}(\phi) \quad \text{s.t.} \quad \bar{D}_{\text{KL}}(\pi_{\text{old}} \parallel \pi_{\phi}) \leq \delta. \quad (7)$$

## Linear-quadratic approximation

$$\max_{\Delta\phi} g^{\top} \Delta\phi \quad \text{subject to} \quad \frac{1}{2} \Delta\phi^{\top} F \Delta\phi \leq \delta, \quad (9)$$

**Issue:** Even though we have selected an optimal step size for Problem (9) via Eq. (10), we are not guaranteed that in (7), we

- actually get a descent in  $L_{\text{TRPO}}(\phi)$ ,
- satisfy the constraint  $\bar{D}_{\text{KL}}(\pi_{\text{old}} \parallel \pi_{\phi}) \leq \delta$ .

**Step length selection via backtracking.** Define a *shrinkage parameter*  $\alpha \in (0, 1)$  and

1. Get  $\Delta\phi$  by solving (9).
2. Calculate candidate  $\phi' = \phi_{\text{old}} + \Delta\phi$  for the next iterate.
3. Check the following two conditions (recall that  $L_{\text{TRPO}}(\phi_{\text{old}}) = 0$ ):

$$(I) \quad L_{\text{TRPO}}(\phi') > 0? \quad \text{and} \quad (II) \quad \bar{D}_{\text{KL}}(\pi_{\text{old}} \parallel \pi_{\phi'}) \leq \delta?$$

**Note:** Condition (II) is checked point-wise on the training batch:

$$\bar{D}_{\text{KL}}(\phi_{\text{old}} \parallel \phi_{\text{new}}) = \frac{1}{N} \sum_{i=1}^N D_{\text{KL}}(\pi_{\phi_{\text{old}}}(\cdot | s_i) \parallel \pi_{\phi_{\text{new}}}(\cdot | s_i)).$$

That's why TRPO does not apply to deterministic policies right away!

4. if (I) and (II)  then  $\phi_{\text{new}} \leftarrow \phi'$  and STOP else:  $\Delta\phi \leftarrow \alpha\Delta\phi$  and back to 2.

# TRPO: Final algorithm

**Input:** Initial policy parameters  $\phi^{(0)}$ , trust-region bound  $\delta$ , backtracking parameter  $\alpha \in (0, 1)$ , acceptance threshold  $\epsilon > 0$ .

For iterations  $k = 0, 1, 2, \dots$  until convergence:

1. **Data collection:** Run the current policy  $\pi_{\phi^{(k)}}$  in the environment to collect a batch of trajectories  $\mathcal{B} = \{(s_t, a_t, r_t)\}$ .
2. **Advantage estimation:** Estimate the advantage values  $A_{\theta^{(k)}}(s, a)$  for all sampled steps.
3. **Compute policy gradient ( $g$ ):** Evaluate the standard policy gradient vector  $g$  using the collected samples:

$$g = \frac{1}{|\mathcal{B}|} \sum_{s, a \in \mathcal{B}} \nabla_{\phi} \log \pi_{\phi}(a|s) \Big|_{\phi=\phi^{(k)}} A_{\phi^{(k)}}(s, a)$$


4. **Compute Step Direction ( $x$ ) via Truncated CG** (10 to 20 iterations) to solve the linear system  $Fx = g$  for  $x \approx F^{-1}g$ .
5. **Maximal step size ( $\beta$ ):** Compute  $\beta = \sqrt{\frac{2\delta}{x^T F x}}$ . Set the full proposed step direction:  $\Delta\phi = \beta x$ .
6. **Backtracking:** Find largest factor  $\gamma = \alpha^j$  (for  $j = 0, 1, 2, \dots$ ) such that  $\phi_{\text{new}} = \phi^{(k)} + \gamma\Delta\phi$  satisfies:

$$L_{\text{TRPO}}(\phi_{\text{new}}) > 0 \quad \text{and} \quad \bar{D}_{\text{KL}}(\phi^{(k)} \parallel \phi_{\text{new}}) \leq \delta$$

7. **Policy Update:** Update parameters to the safe, verified step:  $\phi^{(k+1)} = \phi^{(k)} + \gamma \Delta \phi$  (or *reject* if no  $\gamma$  can be found).

# TRPO: The big picture and relation to actor-critic

## The Big Picture Structure





## Relation to actor-critic

- At its core, TRPO is an Actor-Critic algorithm in practice:
  - It uses a value network critic to estimate  $A_{\theta^{(k)}}(s, a)$ .
- However, it was derived from first principles using policy gradient theory.
  - The primary goal of the paper was to solve a massive mathematical flaw in early actor-critic methods: the lack of monotonic improvement guarantees.



# Example: Half-Cheetah

As an example, we study the **Half Cheetah** example from the MuJoCo library.

- $\mathcal{S}$ : 17 states.
- $\mathcal{A}$ : 6 actions.
- Reward: Forward-Cost - Control-Cost.



Results with TRPO (from the **RL Baselines3 Zoo**)



# Proximal policy optimization (PPO)

# Drawbacks of TRPO

## 1. Large computational overhead


- TRPO relies on the Conjugate Gradient (CG) method to compute the matrix-free step direction  $F^{-1}g$ .
- While CG is much cheaper than explicit matrix inversion, it still requires running an inner loop for every single policy update, evaluating two backpropagation passes per CG step.

## 2. Code complexity and fragility


- TRPO requires custom optimization pipelines. You cannot simply hand a TRPO objective to a standard deep learning optimizer like Adam.
- Requires writing both a dedicated Conjugate Gradient solver and backtracking line search.

## 3. Incompatibility with shared architectures

- In modern deep RL, it is common to use a shared network architecture with two heads for the policy actions (Actor) and the state-value estimates (Critic).
- Because TRPO's constraint is calculated purely using the Fisher Information Matrix of the policy, it has no mathematical way to



Two separate networks



A shared network  
(e.g., a shared trunk  
and separate heads)

Inspired by Sergey Levine's [CS285 lecture](#).

constrain or properly scale the value function parameters.

⇒ we are forced to maintain two entirely separate neural networks.

# Avoiding the constraint

Breakthrough leading to **Proximal Policy Optimization (PPO)** (Schulman et al. 2017):

⇒ no need for an expensive constraint boundary for safe updates.

⇒ we can enforce a trust region inside the objective function by penalizing too large changes in the policy.

## The clipped surrogate objective in PPO

$$L_{\text{CLIP}}(\phi) = \hat{\mathbb{E}}_t [\min(\kappa_t(\phi) A_t, \text{clip}(\kappa_t(\phi), 1 - \epsilon, 1 + \epsilon) A_t)], \quad (11)$$

where we have simplified the notation in several ways:

- $\kappa_t(\phi) = \frac{\pi_\phi(a_t|s_t)}{\pi_{\phi_{\text{old}}}(a_t|s_t)}$  is the *importance sampling* ratio.
- $A_t$ : *estimated advantage* at time  $t$ , i.e.,  $A_\theta(s_t, a_t)$  (typically using a *critic*).
- $\epsilon$  is a hyperparameter (typically 0.1 or 0.2) determining how far the policy is allowed to drift.
- $\hat{\mathbb{E}}_t[f_t] = \frac{1}{T} \sum_{t=1}^T f_t$  denotes an *empirical expectation* (i.e., a sample average) which, in the limit  $T \rightarrow \infty$ , yields our analytical expectation  $\mathbb{E}_{s \sim \rho_{\pi_\phi}, a \sim \pi_\phi}[f_t]$ .

💡 This notation was chosen in (Schulman et al. 2017) for two reasons:

1. It aligns closely with the implementation.

**Recall:** The TRPO optimization problem (Eq. (7)) for comparison:

$$\max_{\phi} \mathbb{E}_{s \sim \rho_{\pi_{\text{old}}}, a \sim \pi_{\text{old}}} [\kappa(\phi) A_{\pi_{\text{old}}}(s, a)]$$

subject to

$$\mathbb{E}_{s \sim \rho_{\pi_{\text{old}}}} [D_{\text{KL}}(\pi_{\text{old}}(\cdot|s) \parallel \pi_\phi(\cdot|s))] \leq \delta.$$

or, in short:

$$\begin{aligned} & \max_{\phi} L_{\text{TRPO}}(\phi) \\ \text{s.t. } & \bar{D}_{\text{KL}}(\pi_{\text{old}} \parallel \pi_\phi) \leq \delta. \end{aligned}$$

⇒ **clipping replaces the constraint!**


2. It handles both infinite and finite-time trajectories.

# Understanding the clipped surrogate objective (1)


$$L_{\text{TRPO}}(\phi) = \mathbb{E}_{s \sim \rho_{\pi_{\text{old}}}, a \sim \pi_{\text{old}}} \left[ \underbrace{\frac{\pi_{\phi}(a|s)}{\pi_{\text{old}}(a|s)}}_{=\kappa(\phi)} A_{\pi_{\text{old}}}(s, a) \right] \quad \text{vs.} \quad L_{\text{CLIP}}(\phi) = \hat{\mathbb{E}}_t [\min(\kappa_t(\phi) A_t, \text{clip}(\kappa_t(\phi), 1 - \epsilon, 1 + \epsilon) A_t)].$$

The clipping function  $\text{clip}(r_t(\phi), 1 - \epsilon, 1 + \epsilon)$  bounds the importance sampling ratio  $\kappa_t$  to

- at most  $1 + \epsilon$  if  $A > 0$ ,
- at least  $1 - \epsilon$  if  $A < 0$ .



The **min** function ensures that we take the minimum of the clipped and unclipped objective  $\Rightarrow L_{\text{CLIP}}$  is a lower (i.e., pessimistic) bound on the unclipped objective.



Source: (Schulman et al. 2017)

Loss function behavior for interpolating policy  $\pi_\alpha = \alpha\pi_\phi + (1 - \alpha)\pi_{\text{old}}$ . (Source: (Schulman et al. 2017))

# Understanding the clipped surrogate objective (2)

$$L_{\text{TRPO}}(\phi) = \mathbb{E}_{s \sim \rho_{\pi_{\text{old}}}, a \sim \pi_{\text{old}}} \left[ \underbrace{\frac{\pi_{\phi}(a|s)}{\pi_{\text{old}}(a|s)}}_{=\kappa(\phi)} A_{\pi_{\text{old}}}(s, a) \right] \quad \text{vs.} \quad L_{\text{CLIP}}(\phi) = \hat{\mathbb{E}}_t [\min(\kappa_t(\phi) A_t, \text{clip}(\kappa_t(\phi), 1 - \epsilon, 1 + \epsilon) A_t)].$$

## Examples

### 1. Positive Advantage ( $A_t > 0$ )

- This means the action was good, and the optimizer wants to make it more likely by increasing  $\kappa_t(\phi)$  above 1.0.
- However, the clipping term caps the reward at  $1 + \epsilon$ .
- The **min** operator ensures that even if  $\kappa_t(\phi)$  grows (for instance, to 1.5 or 2.0), the optimizer receives no extra gradient incentive to push it further.
- It flattens the objective landscape, stopping the policy from overshooting.

### 2. Negative Advantage ( $A_t < 0$ )

- This means the action was bad, and the optimizer wants to make it less likely by dropping  $\kappa_t(\phi)$  below 1.0.
- The clipping term caps this at  $1 - \epsilon$ .
- **Important:** if the policy makes a very large change that makes a good action suddenly zero probability ( $\kappa_t \rightarrow 0$ ), the objective drops significantly, keeping the policy in the safe zone.

# PPO: Final algorithm

**Input:** Initial policy parameters  $\phi^{(0)}$ , clipping parameter  $\epsilon$ , gradient descent algorithm (Adam, RMSprop, ...)

For iterations  $k = 0, 1, 2, \dots$  until convergence:

1. **Data collection:** Run the current policy  $\pi_{\phi^{(k)}}$  in the environment to collect a batch of trajectories  $\mathcal{B} = \{(s_t, a_t, r_t)\}$ .
2. **Advantage estimation:** Estimate the advantage values  $A_{\theta^{(k)}}(s, a)$  for all sampled steps.
3. **Optimize surrogate loss:**  $\phi^{(k+1)} = \arg \min_{\phi} L_{\text{clip}}(\phi)$  via mini-batch training.

# TRPO and PPO are on-policy!

- At first glance, TRPO and PPO can look off-policy because their objective functions use importance sampling:

$$\kappa_t(\phi) = \frac{\pi_\phi(a_t|s_t)}{\pi_{\phi_{\text{old}}}(a_t|s_t)}.$$

- However, the intent behind it here is different:
  - In **off-policy RL**:  $\pi_{\theta_{\text{old}}}$  could be a completely different policy from long ago.
  - In **TRPO/PPO**:  $\pi_{\theta_{\text{old}}}$  is the immediate, exact current policy that just finished stepping through the environment with the policy we're trying to update.

## The Mathematical Reason: Local Approximations

- As derived earlier, the surrogate objectives for both TRPO and PPO rely on a critical assumption: the states we use to calculate the loss must be sampled from a distribution that is nearly identical to the new policy's distribution ( $\rho_{\pi_\phi} \approx \rho_{\pi_{\phi_{\text{old}}}}$ ).
- Because their mathematical bounds collapse if the data isn't fresh, TRPO and PPO are **strictly on-policy**. If we tried to feed TRPO or PPO a batch of off-policy data collected by an entirely different policy/network:
  - In TRPO, the true KL divergence would exceed the trust region  $\delta$ , the backtracking line search would reject the step, and the algorithm would freeze.
  - In PPO, the importance ratio  $\kappa_t(\phi)$  would land far outside the  $(1 - \epsilon, 1 + \epsilon)$  safety window, the clipping mechanism would zero out the gradients, likely resulting in no learning at all.

## The one small exception: “Near-policy” (multiple epochs)



If changes are small, then we can try to use data in multiple consecutive iterations. Tradeoff:

— Data is not entirely on-policy.      + More data  $\Rightarrow$  lower variance.



# Example: Half-Cheetah

As an example, we study the **Half Cheetah** example from the MuJoCo library.


- $\mathcal{S}$ : 17 states.
- $\mathcal{A}$ : 6 actions.
- Reward: Forward-Cost - Control-Cost.



Results: TRPO vs. PPO (from the [RL Baselines3 Zoo](#))



TRPO



PPO

# Comparison of the algorithms we have seen in part I

# Comparison of the algorithms we have seen in part I

- Standard policy gradient.
- Natural policy gradient  $\Rightarrow$  Balancing parameters via the Fisher information matrix.
- TRPO  $\Rightarrow$  Maximize surrogate objective subject to KL trust region.
- PPO  $\Rightarrow$  Replace hard trust region with clipped objective.

| Feature                 | Natural Policy Gradient                                                              | TRPO                                                                        | PPO                                                           |
|-------------------------|--------------------------------------------------------------------------------------|-----------------------------------------------------------------------------|---------------------------------------------------------------|
| Core Concept            | Scaled step sizes based on how much $\pi$ changes, not $\phi$                        | Strict boundary (trust region) to guarantee policy improvement              | Mimics trust region via simplified clipped objective function |
| Mathematical Constraint | KL-divergence approximation via Fisher information matrix.                           | Strict: $\bar{D}_{\text{KL}}(\pi_{\phi'} \parallel \pi_{\phi}) \leq \delta$ | Soft constraint via clipped surrogate objective               |
| Optimization Method     | Matrix inversion: $\phi \leftarrow \phi + \alpha F^{-1} \nabla_{\phi} L_{\pi}(\phi)$ | Conjugate Gradient (CG) + line search                                       | Stochastic gradient descent (SGD) or Adam optimizer           |
| Complexity              | Very high (inversion of $F$ )                                                        | High                                                                        | Low (first-order optimization)                                |
| Implementation          | Challenging                                                                          | Very challenging                                                            | Simple                                                        |
| Data Efficiency         | Poor                                                                                 | Moderate                                                                    | High                                                          |

# References

Kakade, Sham M. 2001. “A Natural Policy Gradient.” In *Advances in Neural Information Processing Systems*, edited by T. Dietterich, S. Becker, and Z. Ghahramani. Vol. 14. MIT Press.

[https://proceedings.neurips.cc/paper\\_files/paper/2001/file/4b86abe48d358ecf194c56c69108433e-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2001/file/4b86abe48d358ecf194c56c69108433e-Paper.pdf).

Kakade, Sham, and John Langford. 2002. “Approximately Optimal Approximate Reinforcement Learning.” In *Proceedings of the Nineteenth International Conference on Machine Learning*, 267–74. ICML '02. San Francisco, CA, USA: Morgan Kaufmann Publishers Inc.

Peters, Jan, and Stefan Schaal. 2008. “Natural Actor-Critic.” *Neurocomputing* 71 (7): 1180–90.

Schulman, John, Sergey Levine, Pieter Abbeel, Michael Jordan, and Philipp Moritz. 2015. “Trust Region Policy Optimization.” In *Proceedings of the 32nd International Conference on Machine Learning*, edited by Francis Bach and David Blei, 37:1889–97. Proceedings of Machine Learning Research. Lille, France: PMLR.

Schulman, John, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. 2017. “Proximal Policy Optimization Algorithms.” *arXiv:1707.06347*.